

# Runtime Funtimes

what runtimes give us  
techniques to tackle what they don't

cdleary @ nodevember '14

## **We'll cover:**

What's a runtime?

Techniques to add new capabilities

Random acts of hackery

**Standard disclaimer:** don't speak for my employer

# **JavaScript is the universal runtime**

I know I've heard someone say that before...

# **JS is a General Purpose** Programming Language

Why would you ever need anything else?

**JS is just like LISP**

Are you sure about that?

```
var lispSuperPowers =  
  require('./lisp-dsl-defn');  
  
function greet() {  
  "use lisp"  
  {define  
    salutation  
    {list-ref  
      {list  
        "Hi"  
        "Hello"}  
      {random  
        2}}}  
  {define  
    greet  
    {name}  
    {string-append  
      salutation  
      ", "  
      name}}  
  {greet  
    "Chris"}  
}  
  
exports.greet =  
  lispSuperPowers(greet)
```

(whole module)

Show of hands:  
**who thinks this can work**  
with “stock” node.js  
(no C extensions)

[github.com/cdleary/nodevember-hacks](https://github.com/cdleary/nodevember-hacks)

We'll get back to how that works...

I really came to this

**JS conference** to talk about

**CPU instructions**

(only half serious)

Instructions are divided into  
**syscalls** that ask the OS for stuff  
& **everything else** (e.g. add)



**C program**  
compiler-generated  
executable code

my\_program.exe

function calls  
(e.g. malloc,  
send, recv)



**C runtime**  
syscall-oriented  
or pre-canned  
routines

libc.so, libm.so, libpthread.so, libdl.so

C program that just  
**computes** “fancy” exit codes  
would require **no syscalls**

But that would be super boring




**JS program**  
source text

my\_program.js



"run me!"



**JS runtime**  
syscall-oriented  
or pre-canned  
routines

JIT codegen



generates  
code to run



function calls  
(e.g. Math.random)  
object allocation



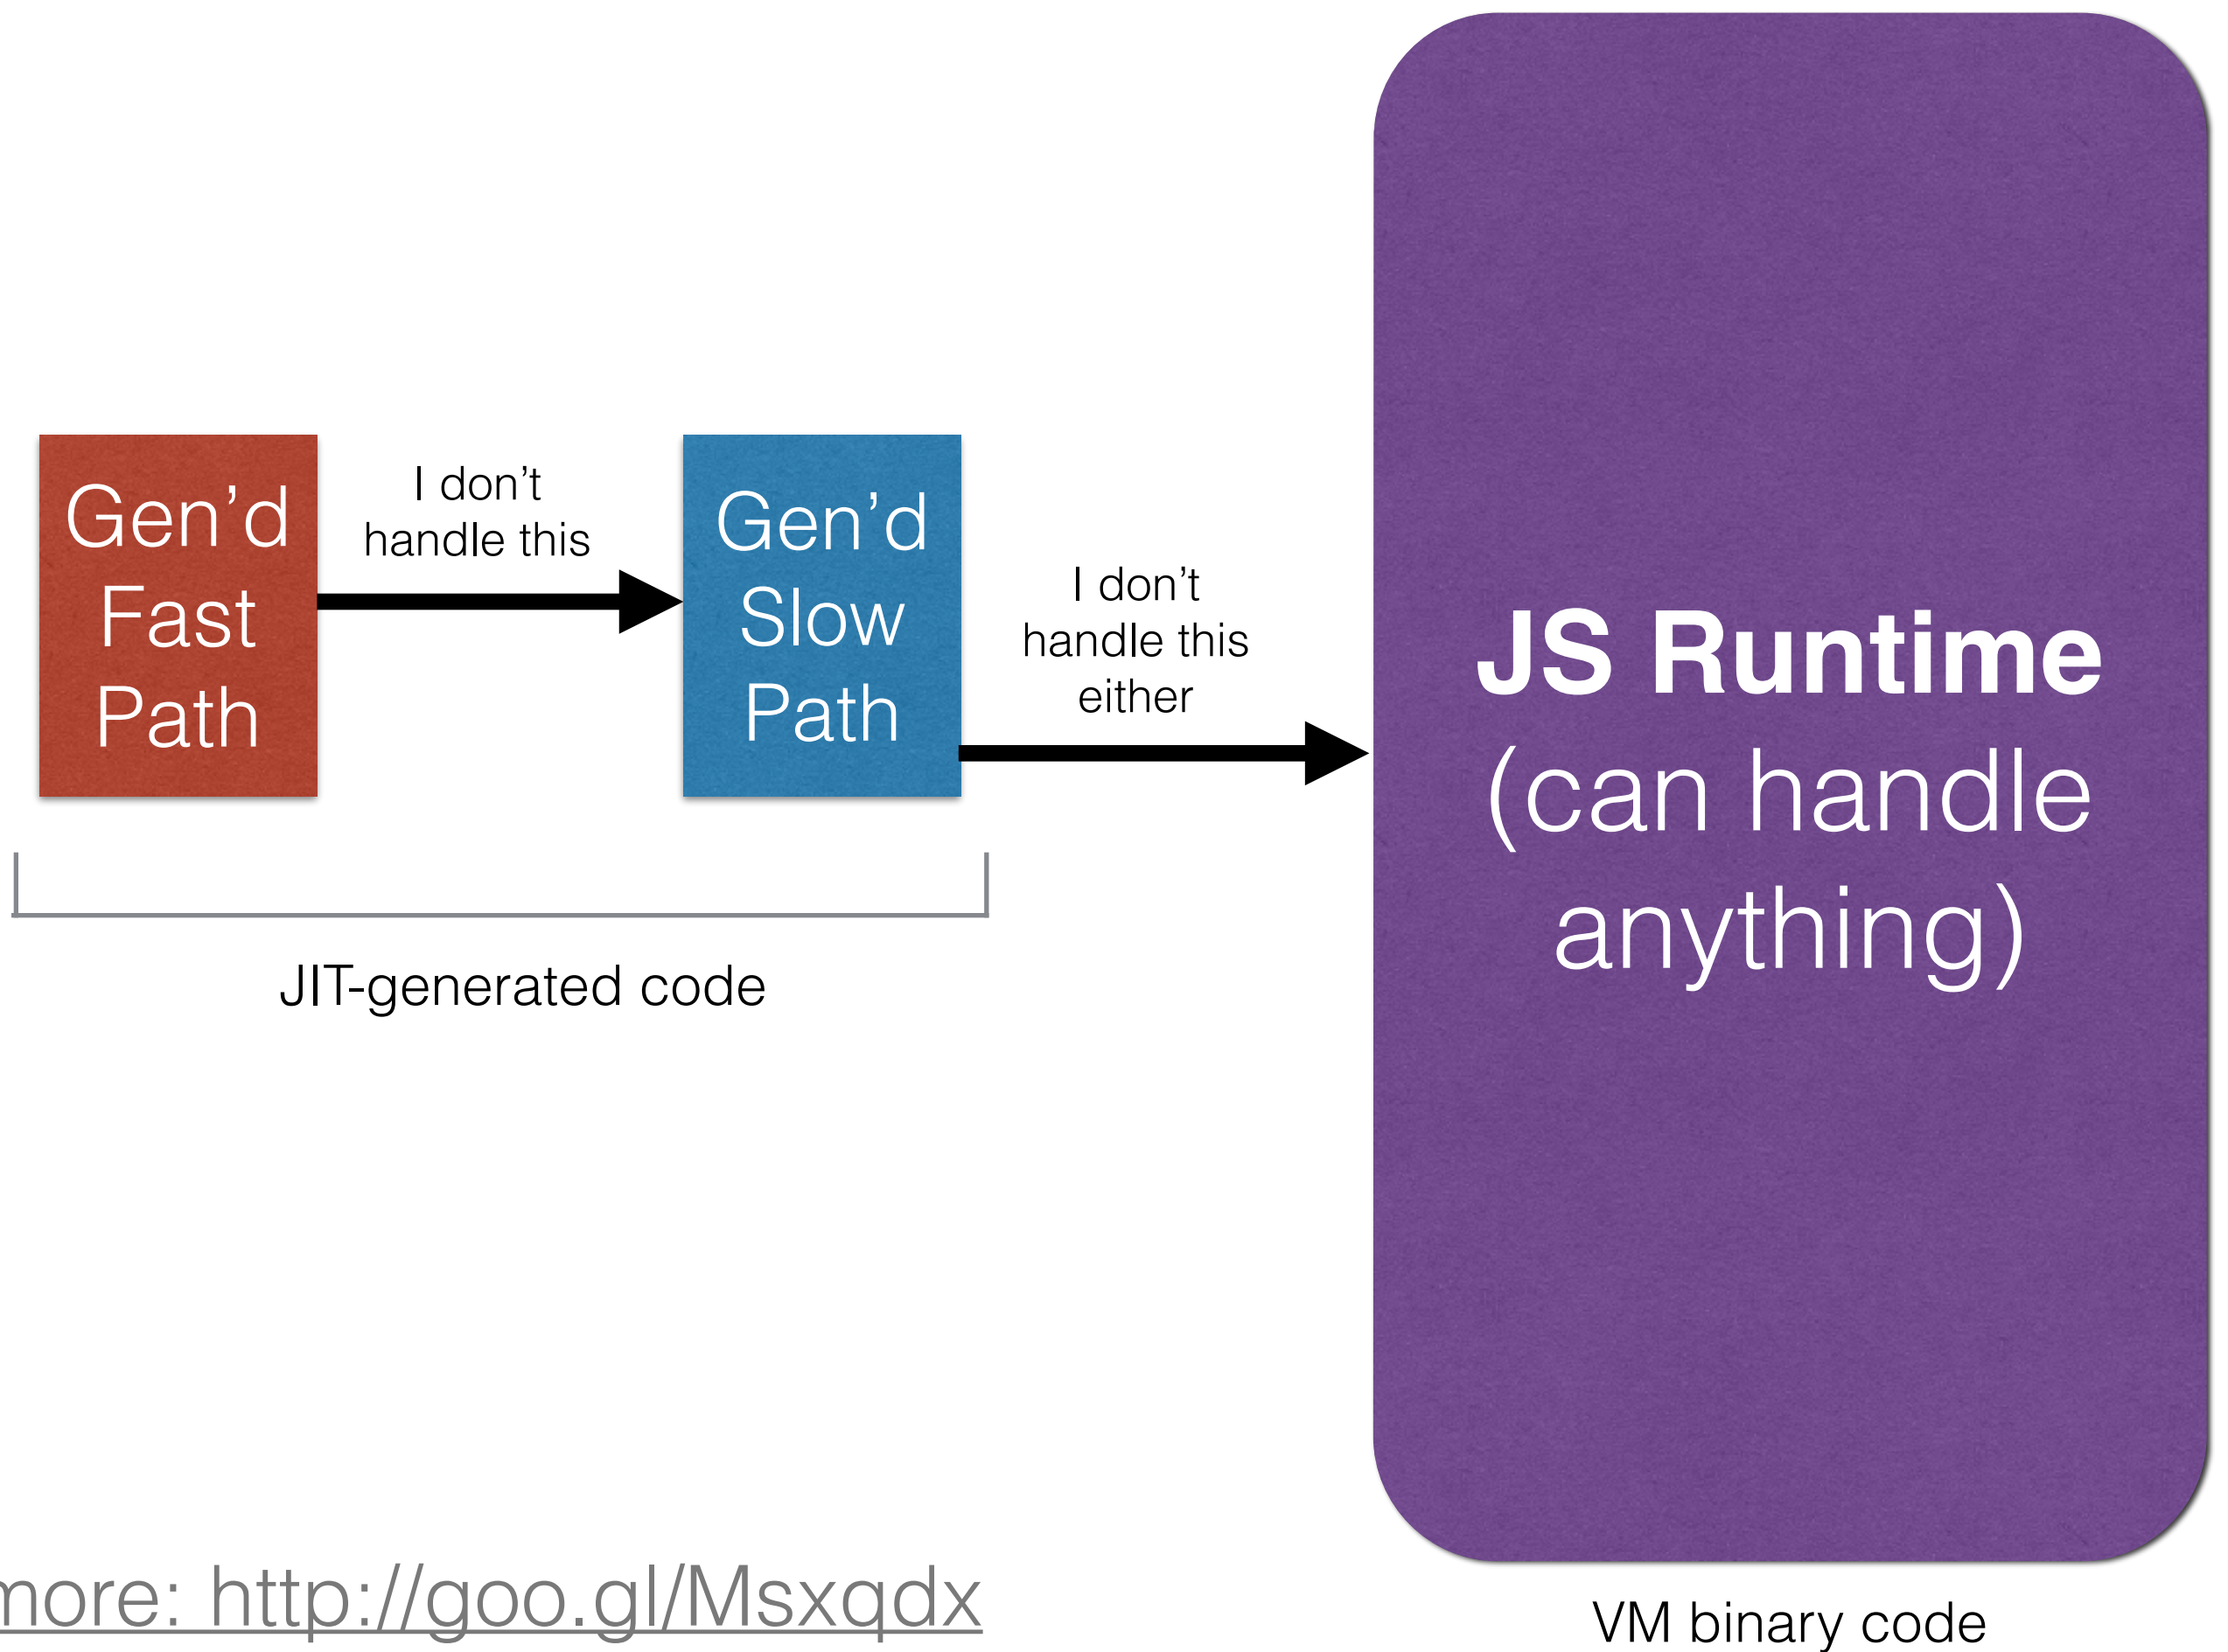
Generated  
executable code

d8.exe

**w/o compilation** step  
(late binding), runtimes  
are **even more critical**

Runtimes are also **fallbacks**  
for **fast path assumptions**  
Violating these results in  
(potentially severe) perf penalties





But did you know?


ES-262 is **useless on its own**

Does not prescribe I/O

Not *even* return codes

The **core** JS runtime is  
**extended** by node  
to make it (async) I/O capable





Runtimes are windows to **the world**  
Program's **execution environment**

I/O, I/O, off to work we go

Runtimes are built to enable  
**semantics of interest**

op a user in our environment wants to perform is **X**  
e.g. addproperty, receive

**node.js runtime is more** than  
V8 with libuv bolted on

defines & supports conceptual constructs  
Streams, EventEmitters, Buffers, ...

Key insight:  
you can **change**  
or **play into** what the **platform**  
you're building on **wants**

“When all you’ve got is a hammer...  
use that hammer to hammer itself into a better  
hammer.”

**–MC Hammer**

Sorry, I forgot,  
he hammered his name into a  
more **condensed, memorable**  
form...

“When all you’ve got is a hammer...  
use that hammer to hammer itself into a better  
hammer.”

**–MC Hammer**



“When all you’ve got is a hammer...  
use that hammer to hammer itself into a better  
hammer.”

~~–MC Hammer~~

–Hammer

(It’s true, check Wikipedia)

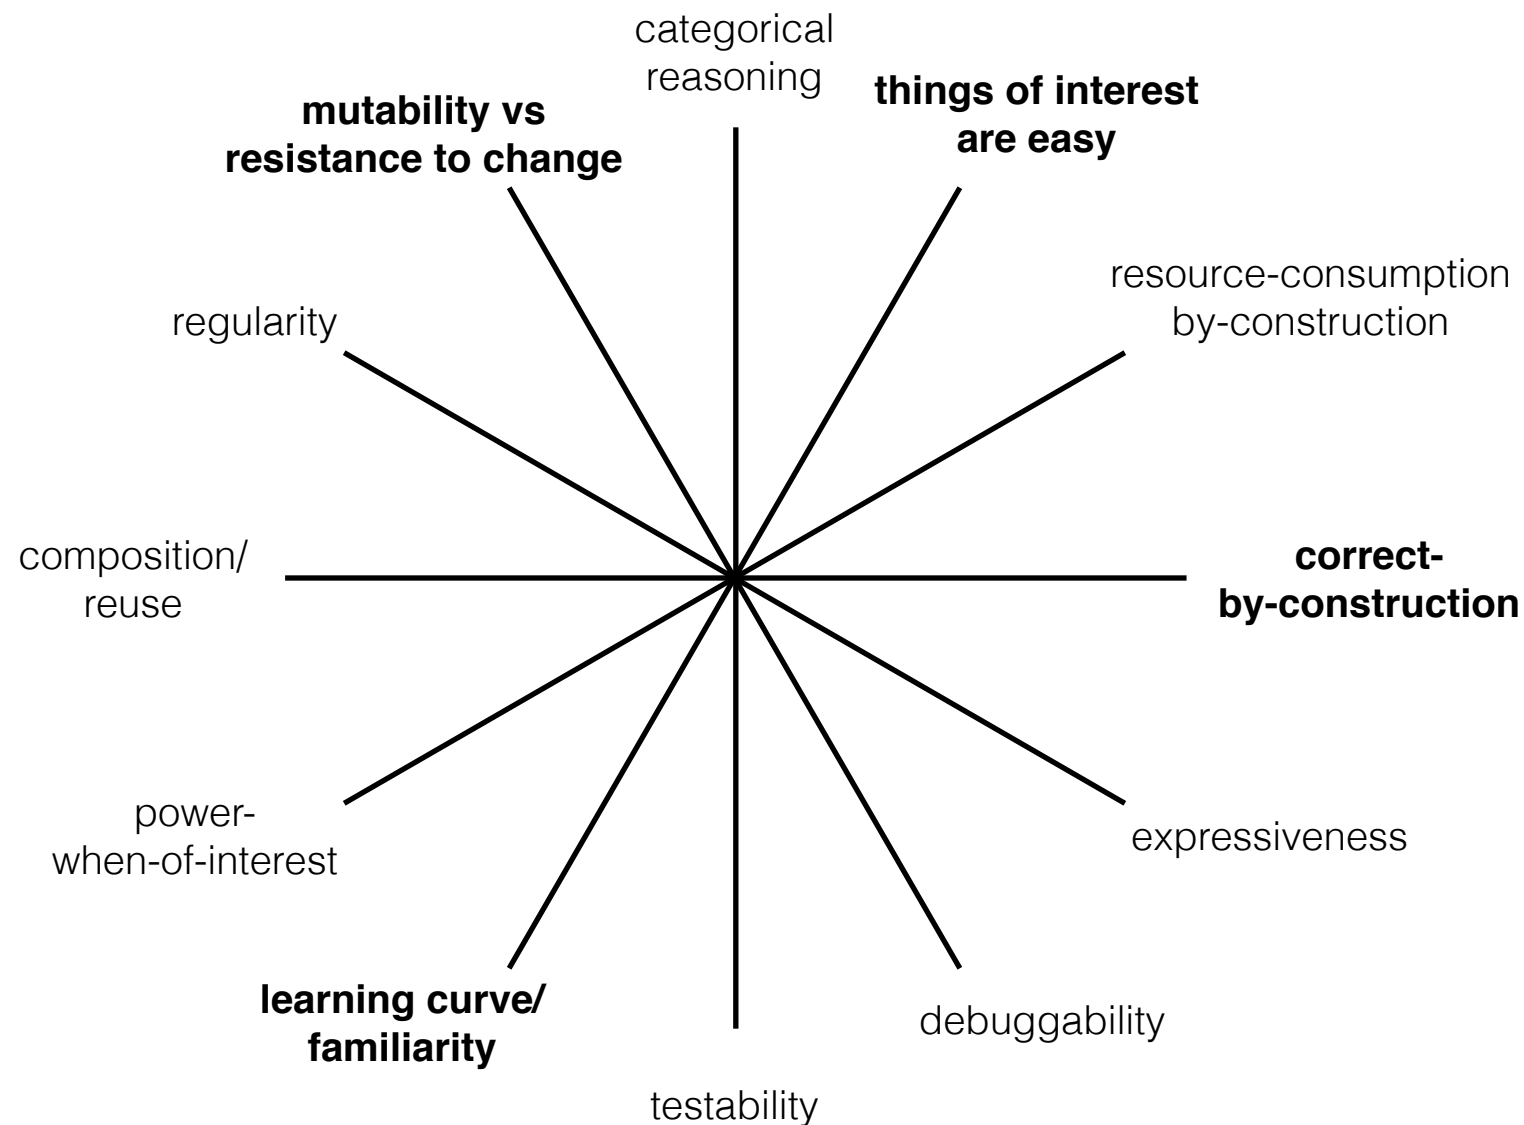
(Except the quote, I totally made that up)

# Technique Rundown

Technique	Approach
Modules	<b>extend lightly</b>
RPCs	<b>bridge</b> , potentially to other platforms
DSLs	<b>constrict</b> , specialize, raise
“Declarative Sandboxes”	<b>constrict w/full language</b> expressiveness
Macros/Desugaring (sweet.js)	<b>embrace, extend</b>
Binary Extensions Hosted FFI Bindings	<b>extend deeply</b> , control tradeoffs
Hosted Natives (JIT Hinting)	<b>symbiotic</b> optimization
Transpilers	<b>reface</b> , change emphasis

**Takeaway:** there's a slew of techniques available to help us achieve our goals

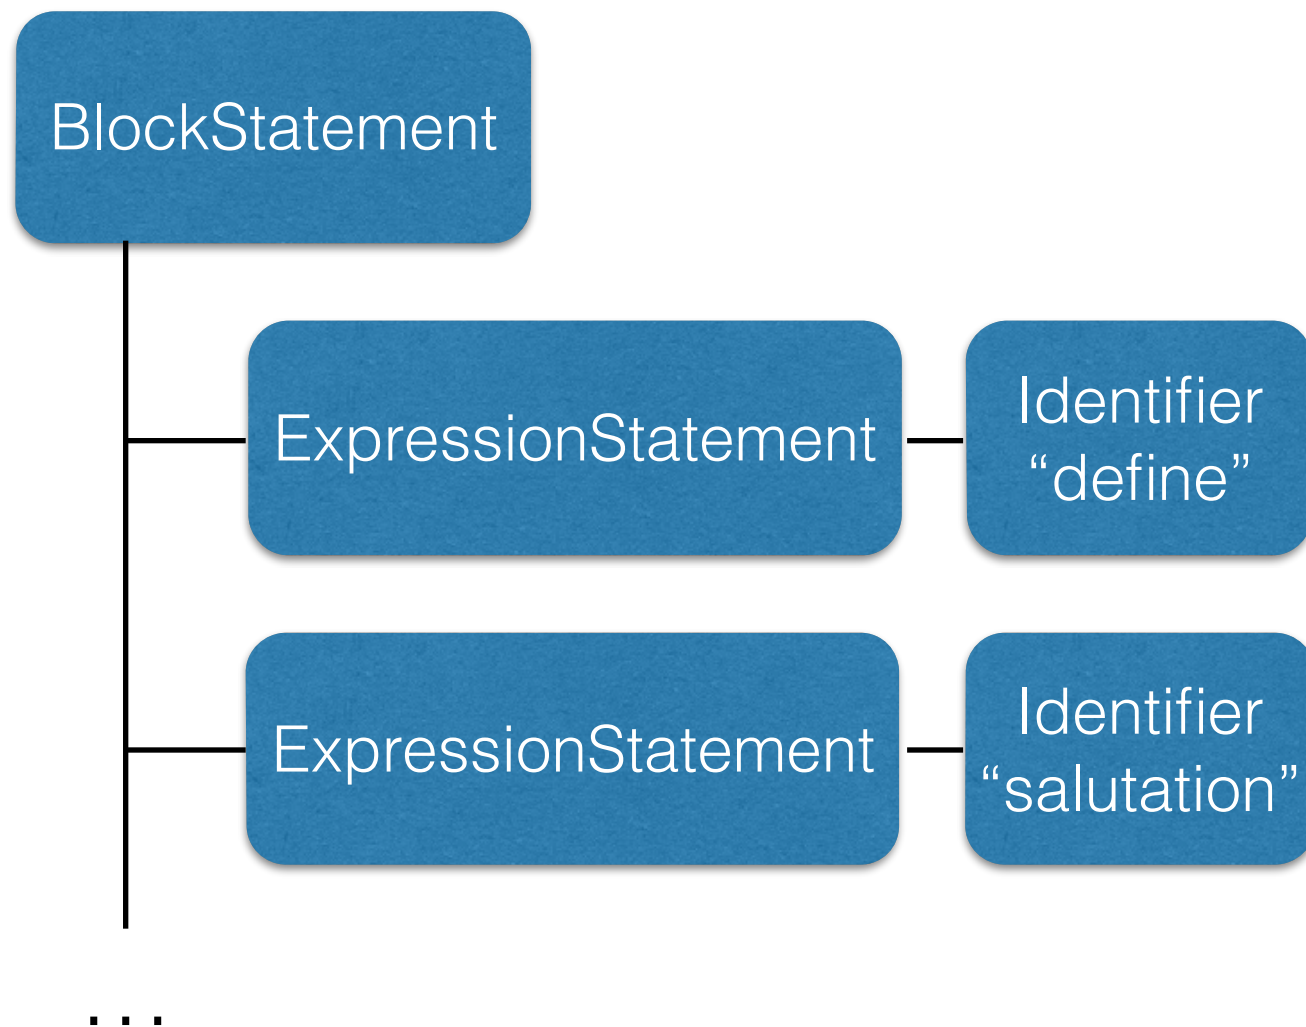
All these techniques are useful because  
**programmability** has **no silver bullet**  
**Therefore,** it's good to have options



# Crux: esprima.parse new Function

```
var lispSuperPowers =  
  require('./lisp-dsl-defn');  
  
function greet() {  
  "use lisp"  
  {define  
    salutation  
    {list-ref  
      {list  
        "Hi"  
        "Hello"}  
      {random  
        2}}}  
    {define  
      greet  
      {name}  
      {string-append  
        salutation  
        ", "  
        name}}  
    {greet  
      "Chris"}  
  }  
  
  exports.greet =  
    lispSuperPowers(greet)
```

# Parse trees



```
{
  "type": "BlockStatement",
  "body": [{
    "type": "ExpressionStatement",
    "expression": {
      "type": "Identifier",
      "name": "define"
    }
  }, {
    "type": "ExpressionStatement",
    "expression": {
      "type": "Identifier",
      "name": "salutation"
    }
  }
],
  ...
}
```

Match against **limited** patterns of JS parse nodes, reusing the JS parser

Either **rewrite** and **new Function** or **interpret** that limited parse tree

**Transform** any syntactically valid JS

Thinking of a use case yet?

µservice:

**do one thing**, do it well, **pipe** together

“I KNOW THIS, IT’S **UNIX!**”



# My dataflow DSL on top of Seneca μservices

seneca.act  
w/continuation

```
var main = senecaPipeline(function runPipeline() {  
  serve&  
  (snapshot2 | snapshot-pair-motion | push-update)*N  
});
```

pipeline  
“snowballs”  
command  
results

async.forever

DSLs could be  
*operator overloading from hell,*  
but provided **semantics** should be  
“whitelist based”

A well done DSL has a very short spec

**function:** short, understandable

**class:** Single Responsibility Principle

**$\mu$ service:** performs a single function

One person's  $\mu$ service  
Is another person's  $\lambda$

More?

write in **sync** DSL

**lower** to **async** program  
across processes in **cluster**?

**Going farther:**  
source maps  
auto-transform hooks  
CPS transform  
partial eval

```
var lispSuperPowers =  
  require('./lisp-dsl-defn');  
  
function greet() {  
  "use lisp"  
  {define  
    salutation  
    {list-ref  
      {list  
        "Hi"  
        "Hello"}  
      {random  
        2}}}  
    {define  
      greet  
      {name}  
      {string-append  
        salutation  
        ", "  
        name}}  
    {greet  
      "Chris"}  
  }  
  
  exports.greet =  
    lispSuperPowers(greet)
```

**Takeaway:** you **don't need a whole compiler** to make your runtime better

**Enabling factor:**  $\mu$ services permit tinkering, are rewrite-friendly: amenable to DSL usage

```
npm install [mynotationhere].js
```

With that one technique covered,  
back to the bigger picture

Runtimes **aren't one-style limited**  
Server-side has a track record of  
doing **very fancy things when**  
**performance really counts**

DSLs w/perf oriented transforms

AoT compilation of Java

Erlang on Xen (look ma! no OS!)

Vortex-style WPO analysis



Runtimes **aren't one-node limited**  
There are **distributed** runtime services

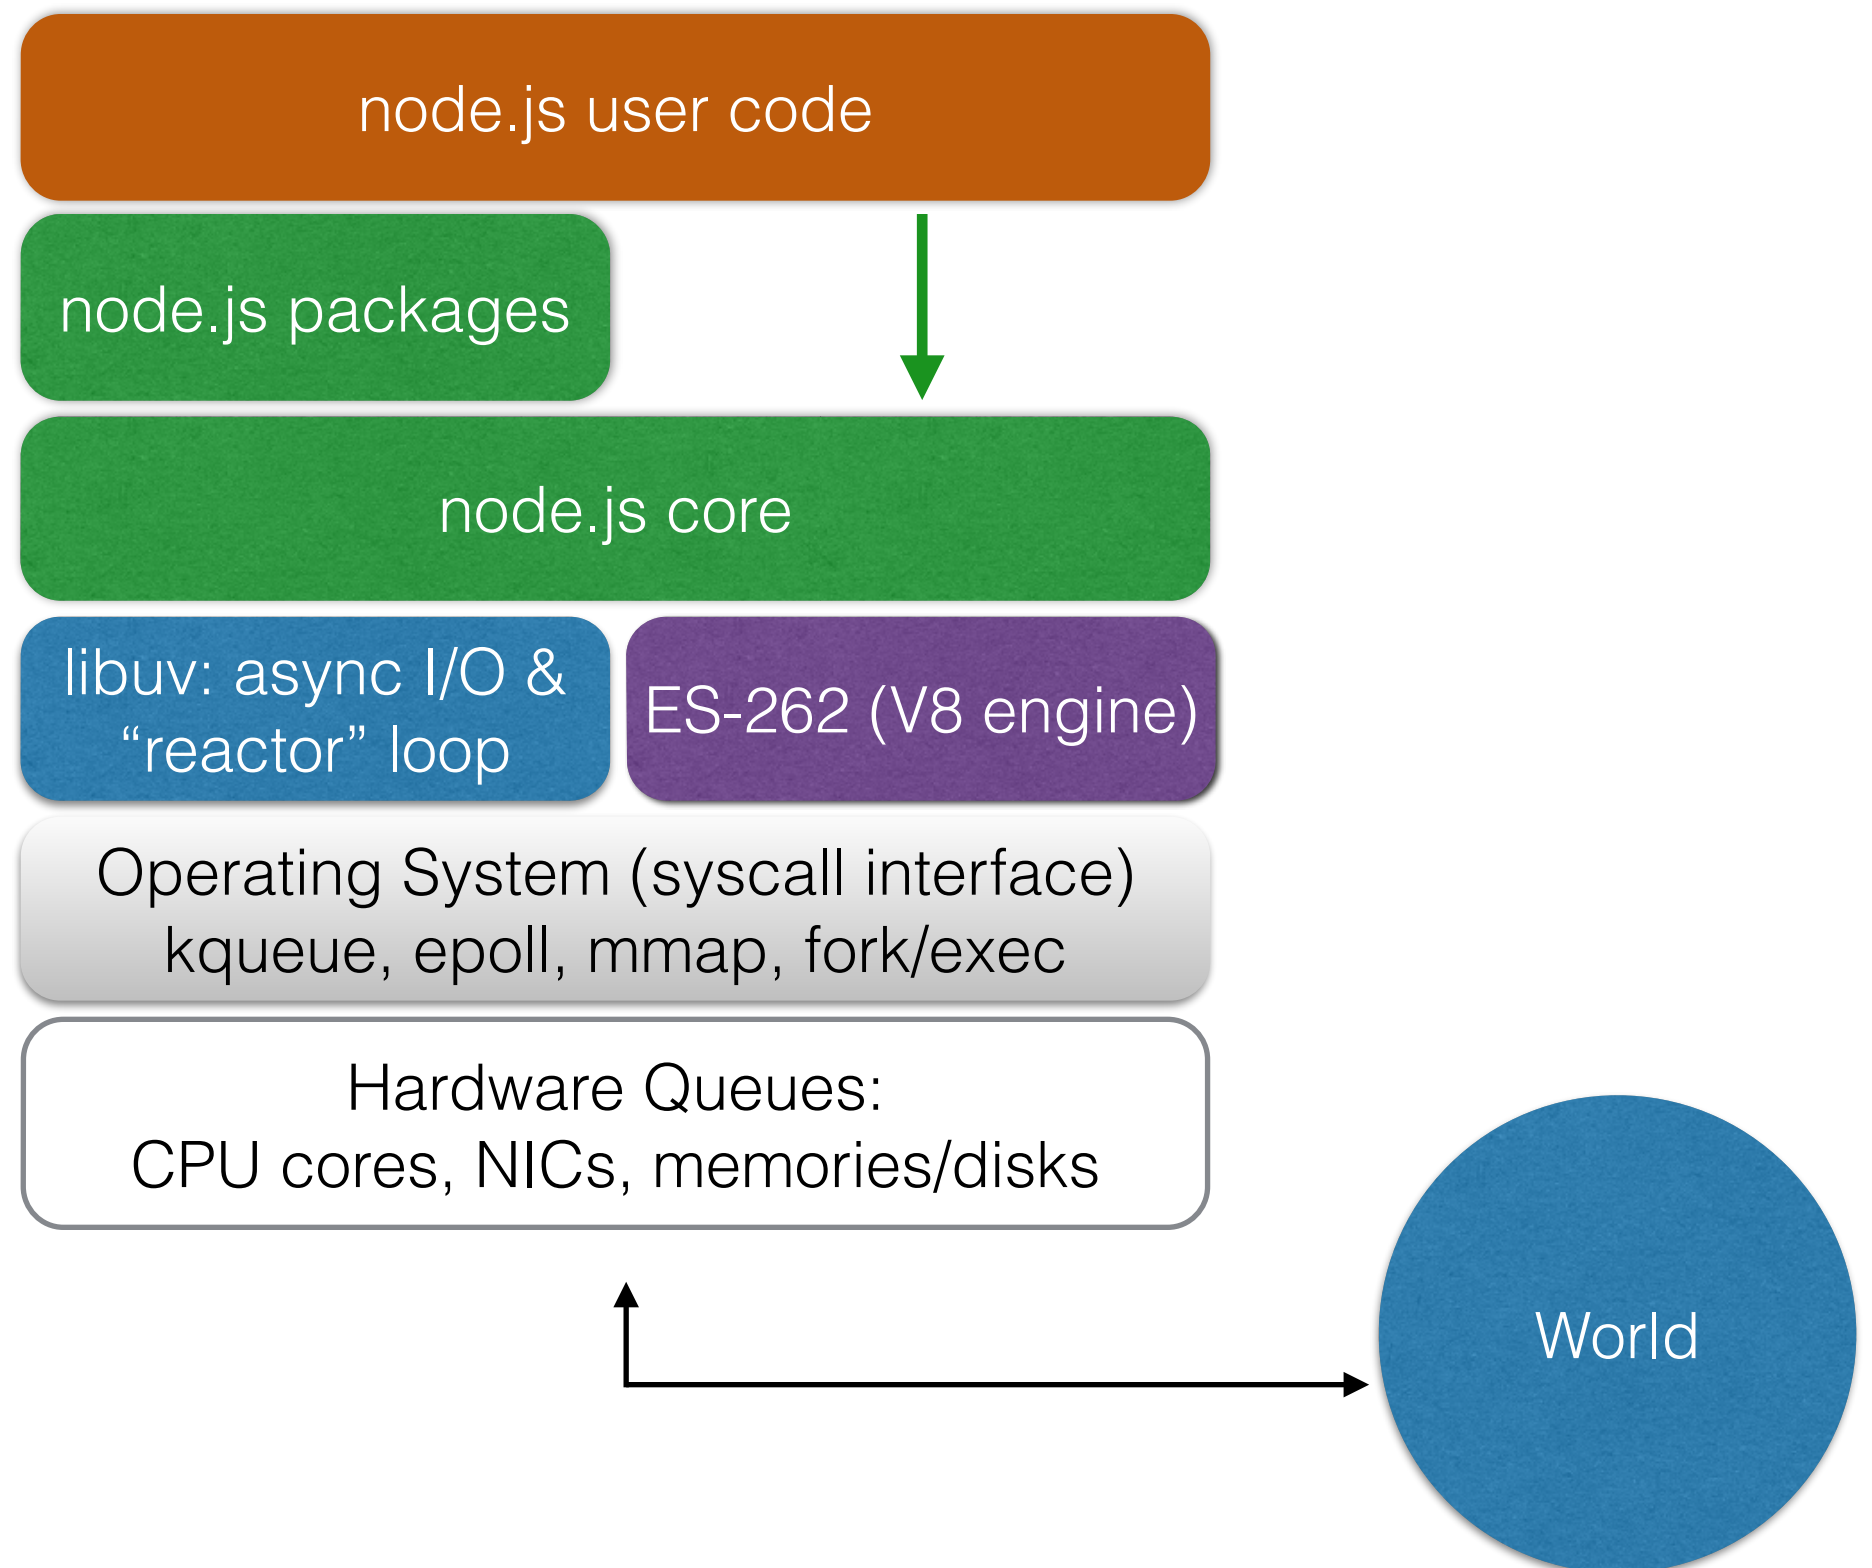
Distributed Smalltalk **cross-node GC** (1987)  
Erlang/BEAM VM **process links** to remote **pids**

Runtime mismatched to use case?

**It can be taught**

In many different ways...

At all the different levels.



Technique	Approach
Modules	extend lightly
RPCs	bridge to other platforms
DSLs	constrict, specialize, raise
“Declarative Sandboxes”	constrict w/full language expressiveness
Macros/Desugaring (sweet.js)	embrace, extend
Binary Extensions Hosted FFI Bindings	extend deeply, control tradeoffs
Hosted Natives (JIT Hinting)	symbiotic optimization
Transpilers	reface, change emphasis

Backup

Thought I was going to talk  
mostly about **maxing perf**



Could only pull **4FPS** from endpoint for  
“my toddler sleep-motion project”, so  
**DSLs** it was :-)

## **JS runtime gives you:**

Aggressive inlining (closures too)

Constant propagation

*Not:* value specialization

*Not:* opt across JS/C boundaries

# Teaching the runtime **more** about performance-critical pieces: “**hosted natives**”

```
// ECMA-262 section 15.5.4.7
function StringIndexOf(pattern /* position */) { // length == 1
    CHECK_OBJECT_COERCIBLE(this, "String.prototype.indexOf");

    var subject = TO_STRING_INLINE(this);
    pattern = TO_STRING_INLINE(pattern);
    var index = 0;
    if (%_ArgumentsLength() > 1) {
        index = %_Arguments(1); // position
        index = TO_INTEGER(index);
        if (index < 0) index = 0;
        if (index > subject.length) index = subject.length;
    }
    return %StringIndexOf(subject, pattern, index);
}
```



runtime implementation is a  
balancing act between  
build **it in**  
build **in it**

# jsctypes/node-**ffi** based “**bolting**”

```
var uv = ctypes.open("libuv.so");  
  var loop = uv.uv_loop_new();  
  var retcode = uv.uv_run(loop);
```

# The dream of Mozilla's **privileged JS runtime** is alive in node (+ ES6)

	node	Moz privileged JS
FFI	node-ffi	jsctypes
Parser Exposure	esprima	Reflect.parse
Message Bus	seneca	XPCOM*
Transpilers?	coffeescript, sweet.js	privileged JS extensions, XUL*

\* Disclaimer: I never worked on Gecko, so this may be completely off.  
In any case, these elements feel familiar.